



Device Context

Apple M5 Pro CPU measured Apple M5 Pro GPU (MLX MPS) measured MTL-FULLER modeled


Latency (ms)




Energy (J)



Avg power (W)



Throughput (img/s)



MPS-relative context: 37.0x lower latency | 1.6x higher energy | 60.4x higher avg power | 37.0x higher throughput

Measured host rows: Apple M5 Pro CPU and Apple M5 Pro GPU (MLX MPS); modeled endpoint: MTL-FULLER. Not benchmark equivalence.